

LLM / / /

: 2026-03-02

: Subagent

202623LLM
24

1.1

pQuant: Decoupled Linear Quantization-Aware Training

- : Wenzheng Zhang, Bingzheng Liu, Yang Hu
- arXiv: 2602.23351
- :
- : scratch QAT

NanoQuant: Efficient Sub-1-Bit Quantization

- : Hyochan Chong, Dongkyu Kim, Changdong Kim, Minseop Choi
- arXiv: 202626
- : LLM 1
- : 1

Quecto-V1: 8-bit Quantized Small Language Models

- : Subrit Dikshit
- arXiv: 2026218
- : 8
- :

TernaryLM: Native 1-Bit Quantization

- : Nisharg Nargund, Priyesh Shukla

- **arXiv:** [2026.02.07](#)
- **QQ:** [QQQQQQQQQQQQQQ1QQQQQQQQQQQQQQQQ](#)
- **QQ:** 1QQQQQQQ + QQQQQQQ

RaBiT: Residual-Aware Binarization Training

- **QQ:** Youngcheon You, Banseok Lee, Minseop Choi [QQ](#)
- **arXiv:** [2026.02.05](#)
- **QQ:** QQQQQQQ LLM QQQQQQQQQQQQQ
- **QQ:** QQQ(± 1) + QQQQQ matmul-free QQ

BPDQ: Bit-Plane Decomposition Quantization

- **QQ:** Junyu Chen, Jungang Li, Jing Xiong [QQ](#)
- **arXiv:** [2026.02.03](#)
- **QQ:** QQ LLM QQQQQQQQQQQQQQQ
- **QQ:** QQQQQQQQQQQQQ

1.2 QQQQQQQQ

QTALe: Quantization-Robust Token-Adaptive Layer Execution

- **QQ:** Kanghyun Noh, Jinheon Choi, Yulhwa Kim
- **arXiv:** 2602.22207
- **QQ:** QQQQQ Token QQQQQQQQQ
- **QQ:** QQQQQQQQQQQQQ

Astro: Activation-guided Structured Regularization

- **QQ:** Xi Chen, Ming Li, Junxi Li [QQ](#)
- **arXiv:** [2026.02.07](#)
- **QQ:** QQQQQQQ LLM QQQQQQQQQQQQQQQQQ
- **QQ:** QQQQQQQQQQQQQ

D²Quant: Accurate Low-bit Post-Training Weight Quantization

- **QQ:** Xianglong Yan, ChengZhu Bao [QQ](#)
- **arXiv:** [2026.02.03](#)
- **QQ:** QQQQQQQQQQQQQ

MatGPTQ: Matryoshka Quantization

- **QQ:** Maximilian Kleinegger, Elvir Crnčević, Dan Alistarh
- **arXiv:** [2026.02.03](#)

- 作者: Matryoshka 数据集
- 作者: 数据集

MoBiQuant: Mixture-of-Bits Quantization

- 作者: Dongwei Wang, Jinhee Kim, Seokho Han
- **arXiv**: 2026.2.21
- 作者: Token 数据集 LLM 数据集
- 作者: token 数据集

Regularized Calibration with Successive Rounding

- 作者: Seohyeon Cha, Huancheng Chen, Dongjun Kim
- **arXiv**: 2026.2.5
- 作者: 数据集

1.3 数据集/数据集

Quantization-Robust LLM Unlearning via LoRA

- 作者: João Vitor Boer Abitante, Joana Meneguzzo Pasquali
- **arXiv**: 2026.2.13
- 作者: 数据集 LLM 数据集
- 作者: 数据集

Differential Privacy + Split Inference

- 作者: Yujie Gu, Richeng Jin, Xiaoyu Ji
- **arXiv**: 2026.2.11
- 作者: 数据集 LLM 数据集
- 作者: 数据集 + 数据集

数据集数据集

2.1 数据集

TOM: Ternary Read-only Memory Accelerator

- 作者: Hongyi Guan, Yijia Zhang, Wenqiang Wang
- **arXiv**: 2026.2.24
- 作者: LLM 数据集
- 作者: 数据集

Compact LLM Deployment in Mobile Edge Computing

- 作者: Ruichen Zhang, Xiaofeng Luo, Jiayi He
- **arXiv:** 2026.02.14
- 主题: 轻量化 LLM 部署
- 关键词: 边缘计算 + 部署

Mapping Gemma3 onto Edge Dataflow Architecture

- 作者: Shouyu Du, MiaoXiang Yu, Zhenyu Xu
- **arXiv:** 2026.02.24
- 主题: Gemma3 在 AMD Ryzen AI NPU 上的部署
- 关键词: 数据流架构 FlowQKV

2.2 模型推理

Scaling State-Space Models with Tensor Parallelism

- 作者: Anurag Dutt, Nimit Shah, Hazem Masarani, Anshul Gandhi
- **arXiv:** 2026.02.24
- 主题: GPU 上的张量并行推理
- 关键词: GPU 推理/加速

SpeContext: Speculative Context Sparsity

- 作者: Jiaming Xu, Jiayi Pan, Hanzhen Wang
- **arXiv:** 2025.11.29
- 主题: 上下文稀疏性
- 关键词: 推理加速

模型推理加速

3.1 模型推理加速

UniComp: Unified Evaluation of LLM Compression

- **arXiv:** 2026.02.11
- 主题: LLM 模型压缩
- 关键词: 模型压缩

SPQ: SVD-Pruning-Quantization Ensemble

- **arXiv:** [20260220](#)
- **链接:** [https://arxiv.org/abs/2602.12201](#)

3.2 模型剪枝

GradPruner: Gradient-Guided Layer Pruning

- **作者:** Wei Huang, Anda Cheng, Yinggui Wang
- **arXiv:** [20260127](#)
- **链接:** [https://arxiv.org/abs/2601.12201](#) LLM 模型剪枝

FineScope: Precision Pruning with SAE

- **arXiv:** [2025](#)
- **链接:** [https://arxiv.org/abs/2508.12201](#) SAE 模型剪枝 LLM 模型

3.3 模型蒸馏

SlimMoE: Expert Slimming and Distillation

- **作者:** Zichong Li, Chen Liang, Zixuan Zhang
- **arXiv:** [20250623](#)
- **链接:** [https://arxiv.org/abs/2506.12201](#) MoE 模型蒸馏

daDPO: Distribution-Aware DPO

- **作者:** Zhengze Zhang, Shiqi Wang, Yiqun Shen
- **arXiv:** [20250623](#)
- **链接:** [https://arxiv.org/abs/2506.12201](#) DPO

EPiC: Edge-Preserving CoT Condensation

- **作者:** Jinghan Jia, Hadi Reisizadeh, Chongyu Fan
- **arXiv:** [20250604](#)
- **链接:** [https://arxiv.org/abs/2506.01220](#)

模型量化-推理-部署

QVLA: Vision-Language-Action Model Quantization

- **作者:** Yuhao Xu, Yantai Yang, Zhenyang Fan
- **arXiv:** [20260203](#)
- **链接:** [https://arxiv.org/abs/2602.01220](#) VLA 模型量化

- 标题: 时空感知量化

AdaTSQ: Temporal-Sensitivity Quantization for DiT

- 作者: Shaoqiu Zhang, Zizhong Ding, Kaicheng Yang
- arXiv: 2026.02.10
- 模型: Diffusion Transformer 时空感知
- 链接: 论文/代码

时空感知量化

Price of Universality in Vector Quantization

- 作者: Alina Harbuzova, Or Ordentlich, Yury Polyanskiy
- arXiv: 2026.02.05
- 模型: 通用向量量化 0.11 精度
- 链接: 论文/代码

Rethinking Perplexity: Input Length Impact

- 作者: Letian Cheng, Junyan Wang, Yan Gao
- arXiv: 2026.02.03
- 模型: 重新思考 perplexity LLM 输入长度影响

重新思考 perplexity

6.1 量化技术

1. 量化技术: 1-bit/2-bit 量化技术 RaBiT TernaryLM 量化技术
2. Token 量化: MoBiQuant QTALe 量化 token 量化技术
3. 量化技术: 量化技术 Gemma3 NPU 量化技术
4. 量化技术: Astro Quantization-Robust Unlearning 量化技术
5. 量化技术: VLA 量化 Diffusion Transformer 量化技术

6.2 量化技术

模型	量化技术	精度	链接
1-bit 模型	TernaryLM, RaBiT	精度	链接

芯片	方案	精度	芯片
ARMv8	MoBiQuant	16bit	ARMv8
ARMv8	D ² Quant, MatGPTQ	16bit	ARMv8
ARMv8	TOM, Gemma3 NPU	16bit	ARMv8

6.3 模型加速

- 使用 **1-bit** 的稀疏矩阵乘法
- **Token** 级别的稀疏化
- **MoE** 的稀疏化
- 其他加速技术

参考文献

1. pQuant: Towards Effective Low-Bit Language Models via Decoupled Linear Quantization-Aware Training (arXiv:2602.23351)
2. Scaling State-Space Models on Multiple GPUs with Tensor Parallelism (arXiv:2026)
3. TOM: A Ternary Read-only Memory Accelerator for LLM-powered Edge Intelligence (arXiv:2026)
4. MoBiQuant: Mixture-of-Bits Quantization for Token-Adaptive Elastic LLMs (arXiv:2026)
5. SPQ: An Ensemble Technique for Large Language Model Compression (arXiv:2026)
6. Quecto-V1: Empirical Analysis of 8-bit Quantized Small Language Models (arXiv:2026)
7. Compact LLM Deployment and World Model Assisted Offloading in Mobile Edge Computing (arXiv:2026)
8. Quantization-Robust LLM Unlearning via Low-Rank Adaptation (arXiv:2026)
9. QTALE: Quantization-Robust Token-Adaptive Layer Execution for LLMs (arXiv:2602.22207)
10. NanoQuant: Efficient Sub-1-Bit Quantization of Large Language Models (arXiv:2026)
11. Astro: Activation-guided Structured Regularization for Outlier-Robust LLM Quantization (arXiv:2026)
12. TernaryLM: Memory-Efficient Language Modeling via Native 1-Bit Quantization (arXiv:2026)

13. MatGPTQ: Accurate and Efficient Post-Training Matryoshka Quantization (arXiv:2026)
14. D²Quant: Accurate Low-bit Post-Training Weight Quantization for LLMs (arXiv:2026)
15. RaBiT: Residual-Aware Binarization Training for Accurate and Efficient LLMs (arXiv:2026)
16. Mapping Gemma3 onto an Edge Dataflow Architecture (arXiv:2026)
17. Regularized Calibration with Successive Rounding for Post-Training Quantization (arXiv:2026)
18. BPDQ: Bit-Plane Decomposition Quantization on a Variable Grid (arXiv:2026)
19. QVLA: Not All Channels Are Equal in Vision-Language-Action Model's Quantization (arXiv:2026)
20. UniComp: A Unified Evaluation of LLM Compression via Pruning, Quantization and Distillation (arXiv:2026)
21. GradPruner: Gradient-Guided Layer Pruning for LLMs (arXiv:2026)
22. Price of universality in vector quantization is at most 0.11 bit (arXiv:2026)
23. Differentially Private and Communication Efficient LLM Split Inference (arXiv:2026)
24. KBVQ-MoE: KLT-guided SVD with Bias-Corrected Vector Quantization for MoE (arXiv:2026)